

Selling Higgsfield to the enterprise

Land through the agencies and studios that already use us, expand into the CPG brands they serve.

Part 0 — Sharp Start	buyer · pain · riskiest assumption
Part 1 — Strategy & Craft	beachhead · rivals · product POV · outreach · pricing
Part 2 — The Model Landscape	five models · hands-on · multi-model
Part 3 — GPU & Unit Economics	cost of a 30-s ad · infra · objections
Part 4 — Competitive Intelligence	Runway · Adobe
Part 5 — Surprise us	conquest map · expert check

Part 0 — Sharp Start

Buyer. The head of content production at a CPG brand — one level below the CMO, owning the content budget across the in-house studio and its agencies. We often land first through the independent studio that produces for that brand.

Champion: the content lead or designer.

Pain. Consumer tools offer isolated point solutions.

Higgsfield runs the end-to-end creative production pipeline, with the right model at each stage — and breaks the old trade-off: quality, speed and low cost, no longer pick two.

Riskiest assumption. Higgsfield users and ambassadors already work inside the target enterprise accounts.

Part 1 — Strategy & Craft

1.1 — Beachhead vertical

Beachhead: CPG brands — land through independent studios, expand into the brand's in-house studio

The workflow we unlock: one approved creative → hundreds of on-brand variants. A studio takes the brand's approved key visual and packshot and turns it into motion and format variants for every channel and market — in days vs. weeks.

The budget already exists. Brands spend ~\$19 on making content for every ~\$31 on media, and today it goes to studios and agencies ([WFA, 2025](#); [Gartner CMO Spend, 2026](#)). CPG is #1 in US digital video — \$16.9B in 2026, ~2× retail ([IAB, 2026](#)) — and puts ~35% of digital into social, more than any other industry ([eMarketer, 2025](#)). "You can't sell toothpaste through search."

Why now

- Touchpoints went from 5 to 15+ in a decade, and each needs its own content ([BCG, Sept 2026](#)).
- Video buyers using GenAI make or finish 24% of their video ads with it in 2025, 33% in 2026 and 43% in 2027 ([IAB, 2026](#)).
- Yet only 13% of CPG marketers have GenAI in their workflows, and ~60% are bringing content creation in-house ([BCG, Feb 2026](#)); ~75% of CPG companies are stuck in pilots ([BCG & Consumer Goods Forum, Jun 2026](#)).

Why us: the people who make CPG content already use Higgsfield — we turn that into a system for the brand (see 1.2).

1.2 — The two who beat Higgsfield to the room

The two in the room

Why these two: CPG brands and the agencies that produce for them ask for Adobe Firefly and Runway by name in their job ads — Firefly at 28% of 50 CPG brands, Runway at 80% of 20 agency networks (Part 5).

- **Adobe — already on every desk.** Creative Cloud is the default for creative teams, Firefly comes with it, and 20,000+ enterprises rely on Adobe ([Adobe Q3 FY26](#)).
- **Runway — the AI-native standard.** Its focus is studios: film, agencies, premium production — Lionsgate, Omnicom, WPP ([Lionsgate](#); [Runway API](#); [Runway London](#)).

How we win with zero logos: your people already use us

~1M paying customers ([Higgsfield, Sept 2026](#)); 85% of use comes from marketers ([Higgsfield, Jan 2026](#)); business is over half of revenue ([CEO on 20VC, Sept 2026](#)).

Our value: we turn that into a system for your business — built into your process and your agencies' work.

1. **Nothing new to roll out.** Your teams and studios already work in it; nobody needs retraining.
2. **A production stack, not just models.** 1,500+ ready ad templates and 250+ VFX presets turn 50+ models into finished ads ([Higgsfield Enterprise](#)). The Supercomputer agent picks the right model for each step — quality, speed and low cost, no longer pick two — and already runs 20M+ generations a month ([Higgsfield, Aug 2026](#)).
3. **An expert team that knows how to make viral creatives** 150+ in-house creatives — almost half the company — make our own launch films ([CEO on 20VC, Sept 2026](#)). A Forward Deployed Creative sits inside your team ([CEO, Aug 2026](#)), backed by our Academy (400K+ course visitors) ([Higgsfield, Aug 2026](#)).

1.3 — Product POV

Ready for the people who make content — not yet for the business that buys it

3 capabilities that matter — from using Higgsfield and Runway side by side

1. **Every model, priced before it runs.** I see the cost before I press run (Seedance 2.5: 56 credits, Veo 3.1: 80; see 2B). A team knows what a brief costs and swaps the model, not the vendor — OpenAI shut Sora's API on Sept 24 ([OpenAI](#)).
2. **Works natively inside Claude and ChatGPT.** I connected it through MCP and CLI; it runs straight from the chat, no workarounds. Plus plugins for Premiere, Photoshop, DaVinci and Figma ([Higgsfield Enterprise](#)).
3. **Recipes you can reuse.** See an ad or effect you like, open it and run it again on your own product: 1,500+ ad templates and 250+ VFX presets ([Higgsfield Enterprise](#)). In Runway I couldn't find an equivalent library.

3 gaps that will lose deals

1. **Too much on one screen.** Every model and tool sits side by side; it's hard to know where to click, and Runway felt more natural to me. G2 users agree: "make the learning process more intuitive", "tools do not tie with each other" ([G2](#)).
2. **Tools, not outcomes.** We speak hype — viral effects, "a full AI creative team in one chat" ([Supercomputer](#)) — not business. The Enterprise page lists models and capacity; there is no CPG use case with a business metric and no published ROAS or cost-per-asset proof. Runway sells "product feed → ads → results" ([Runway Ads](#)); Adobe sells the content supply chain ([GenStudio](#)).
3. **No brand model.** Our own models handle images and edits, and Soul ID trains characters. But a brand can't train a model on its own asset library — Adobe already does it for Coca-Cola with Firefly Foundry ([Coca-Cola job ad, Aug 2026](#)), and Runway licenses its model weights ([Runway](#)).

Fix first: one guided customer job

1.4 — The asset

Haseeb Kamal Siddiqui, Reckitt Middle East

Haseeb Kamal Siddiqui, Regional Marketing Director, Reckitt Middle East (Dubai, since Dec 2024; [LinkedIn](#)). Reckitt Middle East is building an in-house AI content engine for Dettol and Finish: one proven asset multiplied into hundreds of variants for KSA, UAE and the GCC. The September job ad for it names Higgsfield first among the tools a hire should know ([job ad, Sept 2026](#)).

Then expand: Teresa Fusaro, Senior Director Gen AI Creativity, Global ([RetailX](#)). The same ad says the next job is "taking it to other markets".

Cold email

Subject: Your Dubai content engine, built in Higgsfield

Hi Haseeb,

Your AI Content & Workflow ad lists Higgsfield first. Reading it, it describes Higgsfield Enterprise almost line by line:

- **Workflow where tools are nodes** → Canvas and an agent that route each step to the right one of 50+ models; new models land at launch.
- **Brand rules, character and pack references** → Character Studio and shared brand assets in one workspace, with approvals.
- **One asset → hundreds of variants** → batch and API generation, reframes for every platform, Arabic dubbing with lip-sync.
- **One framework across brands and markets** → a sub-workspace per brand and market, one credit pool.
- **Playbooks for the next market** → a Forward Deployed Creative and an AI Educator on your team.

Could we meet for an hour? We'll go through your brief in detail and demonstrate pipeline for you in Higgsfield.

Vladimir

Enterprise, Higgsfield

LinkedIn DM

Hi Haseeb — exciting to see your deep interest in tech; same here! Your AI content engine for Dettol and Finish proves it. We recently launched Higgsfield Enterprise, and most of that pipeline already lives in it. Coffee in Dubai next week? I'd value your view, and I'm happy to share our insights.

Target list of 10

CPG brands — Expand (buyers one level below the CMO)

#	Company · person	Why them · why now
1	Reckitt · Haseeb Kamal Siddiqui, Regional Marketing Director ME → then Teresa Fusaro, Senior Director Gen AI Creativity, Global	In-house AI content engine in Dubai for Dettol and Finish · its job ad (Sept 2026) names Higgsfield first; Fusaro scales GenAI content across markets (talk, Oct 6, 2026)
2	Amorepacific · Jungwon Heo, SVP, Head of Creative Center	The in-house Creative Center built an AI brand-visual workflow for LABO-H on Higgsfield and Soul (2026) · HERA's job ad lists Higgsfield (Sept 2026)
3	Estée Lauder Companies · Claudine Case Labadie, Global Creative Operations, M·A·C — warm intro through my contacts in ELC marketing	Runs creative operations for ELC's largest makeup brand · new CDMO Aude Gandon (ex-Nestlé) is rebuilding ELC's digital and marketing engine
4	Nestlé · Alexis Herail, Global Head of Content	Runs Nestlé's content studios and content technology · two of its production partners already make Nestlé ads with AI, one on Higgsfield (Sept 2026)
5	AB InBev · Lottie White, Head of Production, draftLine Europe	draftLine is AB InBev's in-house agency · Cannes Creative Marketer of the Year for the third time (Jun 2026); Heckler (row 8) shot its Stella Artois spot in Aug 2026
6	P&G · Paul Chick, Head of Content Production, North America	Leads production across TV, digital, social and e-commerce · Pantene Europe got "five times the assets, five times faster, at one-fifth the cost" with AI (Jul 2026)

Independent studios — Land (and the brand each one opens)

#	Studio · person	Why them · why now
7	Steer Media (Dubai / Karachi) · Arshmaan Ali Khuwaja, Creative Director → Reckitt	AI ad films for Dettol and Finish · the shortest path into Reckitt's Dubai content engine
8	Heckler (Amsterdam, Gulf) · Charu Menon, Managing Director & EP → AB InBev	VFX for Stella Artois "The Gentleman" (Aug 2026) · independent, opened its Gulf office in June 2026
9	Myth Studio / Labs (London) · James Finlay, founder → Nestlé	Hybrid AI animation for Nestlé Compleat (2026) on Runway, Veo and Kling · its AI policy asks for enterprise access and no training — our terms
10	Levector (Ecuador) · David Pimentel, partner → Nestlé	Made Nestlé's Tango campaign on Higgsfield (Sept 2026) · already a self-serve user — the cleanest upgrade to Enterprise

1.5 — The commercial test

Higgsfield Enterprise for Agencies — seats for people, packs for volume

Seats (12-month contract, billed monthly)

Seat	Model access	Parallel jobs	Credits / month	Price / month
Creator	All models and resolutions; unlimited on Higgsfield's own models	2	2,500	\$150
Contributor	Higgsfield's own and open-source models only	1	500	\$75
Reviewer	View, comment, approve	—	—	Free, unlimited

Credit packs on top — shared across the workspace

Pack / month	Per credit	Price / month
100k	\$0.055	\$5,500
300k	\$0.050	\$15,000
500k	\$0.0475	\$23,750
1M	\$0.045	\$45,000

For a 50-person agency: 15 Creators + 10 Contributors (\$3,000) + a 300k pack (\$15,000) = **\$18,000 a month, ≈ \$216K a year**. Reviewers are free.

Assumption: ~340k credits a month (42.5k in seats + a 300k pack) ≈ 50 new single-variant 30-second ads (4,350 credits each), ~170 light versions and 10,000 images (see 3A).

Included: SSO, no training on agency or client content, a workspace per client with usage export, rollover within the year, rates fixed for 12 months, a Forward Deployed Creative on the first live brief.

Why this shape: my starting point is unit economics — Higgsfield's own and open models run at over 80% gross margin, rented frontier models at 20–30% (CEO on 20VC, Sept 2026). The package is built to shift usage to the high-margin models: seats include unlimited use of our own models, frontier models are paid in credits — the customer gets more output for the same spend, and Higgsfield earns more on every seat.

Anchors: Team \$65 / seat with 1,000 credits, Scale \$150 / seat with 2,500 credits; Enterprise: custom credits per seat, unlimited members, per-model discounts (Higgsfield pricing, Oct 6, 2026). Seat roles and pack tiers are my proposal.

Part 2 — The Model Landscape

2A — The map: five models that matter

Read: the quality leaders are China-developed; the cleanest legal paper is Google's. Sell the right model per job and risk profile — not "model X is best".

Model	Wins	Breaks	Sell it for
Seedance 2.5 (ByteDance)	Up to 50 references, 30 s; #1 on dialogue and lip-sync	Priciest (\$0.57/s at 1080p) and slow; not offered in the US	Multi-SKU product and UGC-style ads from brand assets
Wan 3.0 (Alibaba)	#1 on text-to-video and editing; \$0.20/s	No vendor indemnity; 5 jobs at a time	High-volume variants and localisation
MiniMax H3	Cheapest, \$0.08/s; #1 on image-to-video	Studio copyright suit; open-weights licence bars US and EU use	Previz and animatics — not the final asset
Veo 3.1 / Fast (Google)	On Google's indemnity list; C2PA, data residency; Fast \$0.12/s	Lower third on quality boards; 8-s clips; one US region	Legal-reviewed hero work, regulated brands
Kling 3.0 (Kuaishou)	Multi-shot, motion control; the cheapest frontier model per second on Higgsfield	Mid-table quality; prompts processed in China	Product video at volume

- **Why these five:** three quality leaders (Seedance, Wan, H3), the legally safest (Veo), the volume workhorse (Kling).
- **Not in:** Runway Gen-4.5 — #27 on blind voting, not on Higgsfield. Google Omni 1.1 — Preview, outside Google's indemnity.
- **Note:** MiniMax H3 is available through Higgsfield's API only, not the web app.

Comparison matrix (1–5, from Artificial Analysis Elo bands, Oct 6, 2026)

	Seedance 2.5	Wan 3.0	MiniMax H3	Veo 3.1	Kling 3.0
Visual quality	5	5	5	1	2
Motion & physics	5	5	5	1	2
Character consistency	5	5	5	1	1
Controllability	4	5	4	1	2
Speed	2	n/a	3	5	1
Price	1	3	5	2 (Fast: 4)	3

My 2B test agrees: Seedance first, Veo last. Speed: one vendor-run test.

Enterprise pros & cons

	Brand safety	Licence & terms	API & limits	Indemnity
Veo 3.1	+ C2PA, own keys, residency	+ Google Cloud contract	GA; quota or reserved throughput	+ Covered (GA only)
Seedance 2.5	– Studio cease-and-desist (Feb 2026)	– Not offered in the US	10 jobs at a time	Partial
Wan 3.0	Not verified	Outputs labelled where law requires	5 jobs; 99.9% SLA	– Customer indemnifies Alibaba
MiniMax H3	– Studio suit	– Weights barred in US, EU, UK, KR	Not verified	– Licensee indemnifies MiniMax
Kling 3.0	– No provenance statement	+ Commercial use on paid tiers	20 jobs per pack	None found

For the security review: only Google's indemnity is unambiguous; everything else rests on Higgsfield's contract — answer model by model.

Sources: [Artificial Analysis](#), [LMArena](#), [Google indemnified services](#), [Veo 3.1 on Vertex](#), [BytePlus Seedance 2.5](#), [MiniMax H3 licence](#), [Adobe on partner models](#), [OpenAI deprecations](#), [Higgsfield pricing](#).

2B — Hands-on proof

Client & brief. EverySight: adapt its sports ad to outdoor padel in 8 seconds. Keep the athlete, glasses and branding consistent; make the strike and ball flight believable.

Test. Same prompt and the same reference image across four models; each platform used its own image mode. Seedance 2.5 and Veo 3.1 ran through Higgsfield, Gen-4.5 and Kling O3 Pro through Runway. One output each; Kling generated 10 seconds, trimmed to 8.

- **Seedance 2.5:** best overall on action. The end card uses plain type instead of the EverySight wordmark — fix in the edit before delivery.
- **Kling O3 Pro:** closer on movement and ball flight, but contact remains unclear and the glasses wordmark is misspelled.
- **Gen-4.5:** broken ball trajectory and racket interaction; branding is hard to read.
- **Veo 3.1:** broken strike mechanics and ball flight; branding is barely visible.

My pick for this brief: Seedance 2.5. Use it as the production base, then correct the remaining brand details before client delivery. For another client, the pick starts from the brief — what matters most decides the model.

One run per model, 720p; my visual read, frame by frame. Seedance cost 56 credits, Veo 80.

2C — The multi-model argument

First, a note on the procurement call. If I'm arguing "aggregator vs single-model vendor" with procurement, I didn't sell the value earlier. That question is settled with the business owner; procurement negotiates price and terms. What procurement should hear is three things:

1. **Stability.** Many models behind one contract. If a model is withdrawn or degrades — Sora's API was shut on Sept 24 — your teams switch the same day: no new vendor, no new contract.
2. **One window for cost.** One transparent rate card across all models, the price of every generation visible before it runs, one invoice, spend caps per team and per client.
3. **Cheaper overall.** A single-model vendor is hard to negotiate with: your teams use different models, so you can't give any one vendor your full volume. With us the volume is pooled, so the rate falls (see the packs in 1.5) — and you pay flagship prices only where flagship quality is required.

Where it backfires — and the answer

Counter	Answer
Depth vs breadth	Multi-model is how we get depth: the best model for each shot, plus our own layer on top — Soul for images, Soul ID and Character Studio to keep characters and brand faces, Supercomputer to plan the job and pick the models. The result matches or beats any single model on its own
Jack of all trades	We're not a jack of all trades — we get the best out of every model. With a single vendor you get its strengths and all its weaknesses. We take each model's best parts and cover its gaps with another. That is the difference
Dependency on model providers	Depending on one provider is the real risk. We always have another model ready, can test scenarios side by side and switch the same day — that's flexibility and speed, not dependency

Sources: [OpenAI deprecations](#), [Higgsfield Enterprise](#), [Higgsfield pricing](#).

Part 3 — GPU & Unit Economics

3A — What a 30-second ad costs: traditional vs AI

Bottom line: a 30-second paid-social ad costs \$3,803 AI-native against ~\$9K traditional, and takes 6 days instead of 4 weeks. AI removes the shoot: crew, gear, location and actors become ~\$533 of generation.

1. One 30-second ad, two ways (3 variants × 3 formats)

	Traditional	AI-native
Pre-production	\$1,000	\$600
Shoot: director, crew, gear, location	\$5,000	—
Generation	—	\$533
Directing → prompt iteration	in the shoot	\$1,200
Actors and voice	\$1,000	\$450
Post: edit, color, sound, cleanup	\$1,500	\$700
Music and formats	\$500	\$320
Total	≈ \$9,000 (\$8–10K)	\$3,803
Time	4 weeks	6 days

2. Inside the generation line: what drives cost per generation

Cost per usable shot = passes × seconds generated × credits per second × \$ per credit

One 30-second variant: 10 shots × 7 passes on Seedance 2.0 at 1080p (45 credits per 5 s) + 3 product close-ups × 5 extra passes on Veo 3.1 (80 credits per 8 s) = **85 generations, 470 seconds generated to keep 30, 4,350 credits ≈ \$220** at \$0.05 per credit. Three variants that share the close-ups: 3 × 3,150 + 1,200 = 10,650 credits ≈ \$533 — the generation line above.

Driver	Options (Higgsfield rate card, Higgsfield pricing , Oct 6, 2026)	Generation per variant
Model (1080p)	Kling 3.0 · Seedance 2.0 · Veo 3.1	\$28 · \$158 · \$175
Resolution (Seedance 2.0)	720p · 1080p · 4K	\$77 · \$158 · \$385
Clip length per generation	5 s · 8 s · 10 s	\$158 · \$252 · \$315
Passes per shot	3 (bid) · 7 (actual) · 23 (broadcast spot)	\$68 · \$158 · \$518
\$ per credit	Ultra \$0.033 · Enterprise \$0.05 · Team \$0.065	\$104 · \$158 · \$205

One driver at a time, from the body-shot baseline of 70 generations. Passes and resolution move the bill most.

3. ROI: the buyer's math

- **Cost:** \$3,803 vs ~\$9K — **–58%**.
- **Speed:** 6 days vs 4 weeks.
- **Every next variant:** traditionally 20–30% of the first ad, and only if shot the same day; AI-native, ~\$220 of generation plus prompt time.

Stated assumptions: a 30-second ad needs 8–12 shots; bids assume 3 passes per shot, the actual average is 7; 30–90 minutes of human iteration per shot is normal; close-ups need 3–6 hours of cleanup per finished minute; the traditional line split is my estimate within the \$8–10K total.

3B — Infrastructure options

Default: shared cloud. Dedicated capacity only at very large, steady volume. Private VPC or on-prem almost never.

	Shared cloud	Dedicated capacity	Private VPC / on-prem
Cost	Lowest: pay per second, no idle	Higher: you pay for reserved capacity even when it sits idle — 8×H100 ≈ \$18.6K a month (Together)	Highest: GPUs at \$7–12 per hour on a hyperscaler (AWS via Vantage) or your own hardware, plus a team to run it (3C)
Speed & throughput	Same models, same speed — but a shared queue at peaks and vendor caps: Seedance allows 10 jobs at a time per account (BytePlus)	No queue: capacity reserved with the model vendors and on our GPUs	Only what the customer owns; scaling up means finding more GPUs
Data security	Prompts and assets go to model vendors — 40 sub-processors (Trust Center); no training by contract	Same as shared: frontier models still run at their vendors; only Higgsfield's own models get a single tenant	Data never leaves — but only open and Higgsfield models run there: no Veo, Seedance or Kling
Ops burden	Near zero for the customer	Low for the customer, higher for us	Heavy: the customer's team runs it
Upgrades	New models the day they launch	Same for vendor models; our own GPUs need a swap per model	Slowest: every new model is a project, and hardware lives 3–5 years against a two-month model cycle

What I'd tell the buyer: start on shared cloud. It's the cheapest, gets new models first, scales with you and leaves nothing for your team to run. Dedicated capacity pays off only for a very large installation you keep busy around the clock — you pay for the reservation whether you use it or not. Private VPC costs the most, upgrades slowest and can't run the best models — and where would you find GPUs to scale fast?

3C — The objections

"Why don't we just run open-source models on our own GPUs?"

You can — but price the whole thing.

The best models aren't open. Veo, Seedance and Kling run only at their vendors. On your own server you get older open models — Wan 2.2 from 2025, while we already serve Wan 3.0.

The math. One 8×H100 server costs ~\$285K; the current generation, B200, is \$400–500K with a 14–20-week wait. At full load it makes ~600 five-second clips a day. Add power, hosting and the three or four engineers who keep it running — ~\$1.2M a year — and each clip costs ~\$6, or ~\$11 at half load. The same clip in our cloud costs \$0.35–0.44.

It ages fast. A new GPU generation lands every year or two, and new models want it. Your server depreciates against a two-month model cycle, and every model change is a migration project.

So the question isn't whether you can run a model. It's whether running models is your business.

"What happens to our price when GPU costs drop?"

We don't sell GPUs. We sell access to the best models and a result — GPUs are a line deep inside our cost, not yours.

And in practice GPU savings don't arrive: rental is up ~15% this year, and every new flagship launches higher — Seedance 2.5 costs ~50% more than 2.0. A given model does get cheaper over time — Veo fell from \$0.75 to \$0.40 a second — but by then your team wants the next one.

Where you do feel progress is cost per approved ad: newer models need fewer tries per shot, and tries are the biggest driver of the bill. Your bill doesn't shrink; your output per dollar grows.

So your rates are fixed for the year, and once a year we review them together against the market: new models, upstream prices, your volume.

Clip math: server ~\$95K a year over three years + ~\$31K power and hosting + ~\$1.1M for 3–4 engineers ÷ 18,000 clips a month at full load (Wan 2.2, 720p, 5 s); cloud = Higgsfield rate card at \$0.05 per credit (Kling 3.0 / Wan 3.0, 720p, 5 s). Sources: server prices (CloudZero; Mercatus; RackRate), throughput (Wan 2.2), engineer pay (Levels.fyi), rate card (Higgsfield pricing), GPU rental trend (getdeploying), Veo pricing (Google), Seedance prices (Artificial Analysis).

Part 4 — Competitive Intelligence

Competitors: **Runway** and **Adobe (Firefly Enterprise + GenStudio)** — the two from 1.2. All pages accessed Oct 6–7, 2026.

4A — Product & positioning

Claims vs. what users say

	They claim	Users and the vendor's own fine print say	The gap
Runway	"Production-ready AI made for enterprise"; "10x your creative output at 10% the cost" (Enterprise)	"sometimes it takes 10 tries to actually get the result I want" (G2); failed generations still burn credits, which "cannot be automatically reinstated" (Help); the savings figures are "reported by the customers themselves" (AI Media Report); Gen-4.5, launched as top-rated, is now #27 of 48 on blind voting (Arena); its own CRO: "Models are converging" (Runway)	Sells a finished asset; users describe a fast draft machine and pay for every miss
Adobe	"Enterprise-grade content production"; "average 8.5x ROI"; "commercially safe" (Firefly Enterprise)	The #1 pro in TrustRadius's summary of 116 reviews is "accelerates ideation" (TrustRadius); "not quite ready for primetime" (TR); "Firefly model per se is not a great model", partner models are "the best part" (TR); "8.5x" is "up to... over three years" for a modelled company (study); indemnity covers Adobe's own models, Google and OpenAI GA models only with supplemental cover — Runway, Kling, Seedance and beta features are not covered (terms)	Sells safe production at scale; users sketch in Firefly and finish in Photoshop. "Safe" covers the model they rate lowest, not the ones they prefer

Capability matrix — table stakes (others in the deal ship it) · real edge (only this vendor, backed by more than its marketing) · vaporware (the vendor's own page says beta, coming soon or early access)

Capability	Runway	Adobe
Own video model	Table stakes — Gen-4.5 #27 of 48 (Arena); Aleph last, #10 of 10, on video editing (Arena)	Real edge on provenance — trained on licensed content — not on quality (G2)
Editing real footage	Table stakes — Higgsfield's Ad Multiplier does it too; Runway now runs inside Premiere (Runway , Sept 2026)	Real edge — the host apps: Photoshop and Premiere (G2)
Third-party models	Table stakes — Seedance, Kling, Veo on day 0 (Runway)	Table stakes — Seedance and Kling kept off business plans (Adobe)
Agent	Table stakes , early — none of the 10 latest G2 reviews mention it (G2)	Vaporware for enterprise — AI Assistant is "a beta feature" (FAQ)
Enterprise contract	Real edge — IP indemnity outside the liability cap, third-party models included; the customer still checks outputs' legality (§3.2) (terms)	Table stakes — \$3M or 2x fees on enterprise agreements, \$10K on teams licences (pricing)
Brand model	Real edge , narrow — licenses closed model weights (Runway)	Real edge — Firefly Foundry, live at Coca-Cola (job ad)
Loop to ad results	Vaporware — Runway Ads: "Request early access" (Runway Ads)	Table stakes for activation; bulk trafficking vaporware , "coming soon" (GenStudio)

Read: models and agents are table stakes everywhere. The real edges are Runway's contract and Adobe's host apps and brand models. And the two are partly one camp: Gen-4.5 runs inside Firefly ([Business Wire](#), [Dec 2025](#)), Adobe Ventures invested in Runway's Series E ([Runway](#)), Runway plugs into Premiere. Higgsfield's lever is not a feature but the production stack its users already work in (1.2).

Stated vs. observed ICP

	Runway	Adobe
Stated	Enterprise AE: \$1B–5B+ companies, deals \$100K–1M+; mid-market AE: <\$1B, \$40–100K (job post)	"Over 20,000 global enterprises" (Q3 FY26 call); retail, e-commerce and agency use cases (Firefly Enterprise)
Observed — reviewers and installs	G2: 73% small, 7% enterprise (G2); 3,184 installs, 72% in marketing, manufacturing 7.5% (TrustRadius)	G2: 51% small, 21% enterprise (G2); 4,192 installs, manufacturing 11% (TrustRadius)
Observed — who posts	The co-founder of Silverside, an independent AI studio making Coca-Cola ads (LinkedIn); solo creators choosing between Runway and Higgsfield (Reddit)	Designers who finish in Photoshop — "I rarely ship straight out of the app" (G2); an art director finds Adobe's agents "locked to Adobe's own model ecosystem" (Reddit)
Mismatch	Sells top-down to \$1B+ companies; the visible base is creators and small studios	The enterprise business is real but leaves almost no public trail
Against our beachhead	Already at our land stage: independent studios like Silverside build on its API (case)	At our expand stage: the brand's default creative stack

4B — Voice of the market

Base. Adobe — 94 reviews read in full (G2 29, TrustRadius 44, Capterra 21). Runway — 32 (G2 30, Capterra 2); the ≥ 20 bar is met on G2 alone, 22 without the old Gen-2 listing, and TrustRadius has ratings but no text. About 80% of both bases are incentivized, so read themes, not shares. n = reviews raising the theme; one coder.

Adobe

	Theme (n of 94)	Quote
Praise	1. Speed (36)	"Tasks that would have taken hours of manual labor now take seconds" — Sr Marketing Manager, 10,001+ (TR)
	2. Variants at scale (32)	"taking hero content and creating other assets based on it" — Director Marketing, 10,001+ (TR)
	3. Commercially safe (28)	"Image generation is solid and, more importantly, commercially safe" — Sr Analyst, enterprise (G2)
	4. Fits the Adobe stack (27)	"a go-to solution for teams already relying on Adobe tools" — consumer-brand marketer, 10,001+ (TR)
	5. Easy for non-designers (24)	"execute brand ready assets when designers aren't available" (TR)
Complaints	1. Quality (36)	"faces and hands still come out wrong often enough" — COO, mid-market (G2)
	2. Packaging confusion (23)	"More clarity around what is GenStudio and what requires other licenses/products" — Advisor Marketing, 10,001+ (TR)
	3. Cost and credits (22)	"It is expensive, requires cloud subscription" (G2)
	4. Limited control (19)	"fine tune an image once a basic direction has been established without starting over" (TR)
	5. Integration outside Adobe (15)	"Integrating with third-party DAMs still often requires custom development" (TR)

2026 trend: quality is still complaint #1; cost and credits are rising (6 of the 11 reviews since June); "Adobe's own models lag" appears in 5 of 33 Firefly reviews in 2026, none before.

Runway

	Theme (n of 32)	Quote
Praise	1. Idea-to-video speed (19)	"Work that used to take hours of editing can often be finished in minutes" — Social Media Manager, mid-market (G2)
	2. Ease of use (17)	"Runway is very easy and UI/UX is understandable" — Graphic Designer (G2)
	3. Camera and motion control (13)	"quickly get a video with the movement and camera style I want" (G2) — 3 of 13 praise Motion Brush, retired in May 2025

	Theme (n of 32)	Quote
	4. Replaces stock and B-roll (11)	"Instead of spending hours searching through stock footage libraries" (G2)
	5. Many models in one place (9)	"There are many AI models available in one place." (G2, Gen-2 listing)
Complaints	1. Retries, weak prompt adherence (~26)	"sometimes it takes 10 tries to actually get the result I want" — President, small business (G2)
	2. Credit burn; misses still charged (19)	"Even failed generations still use up paid account credits" — Owner, mid-market (G2, Gen-2 listing)
	3. Face and object drift (10)	"Sometimes the subject, face, or background changes between frames" — CGI Lead (G2)
	4. Slow at peak (6)	"Rendering can also be slow during peak hours" (G2)
	5. Learning curve (6)	"some of the more advanced features come with a noticeable learning curve" (G2)

Recency: 12 of the 17 latest G2 reviews raise credits or price. Trustpilot (supplementary, self-selected): Runway 1.1/5 from 354 reviews; 14 of the 20 latest report credits spent with nothing usable, one an agent "consuming all 2,932 credits without asking for my approval" (Trustpilot). Higgsfield scores 4.0, but on a claimed profile that invites reviews, and its own summary mentions generations "stuck on processing, leading to wasted points" (Trustpilot).

YouTube — what surprised or disappointed reviewers

	Video	Timestamp	Takeaway
Adobe	Curious Refuge — 7 platforms, one film (Jan 2026)	02:48	Boards "incredibly intuitive"; credits sold separately; no Chinese models. The same film cost ~\$504 on Firefly, \$150 on Higgsfield, \$72 on Google Flow (results)
Adobe	Dan Kieft — 6 platforms (Nov 2025; has Higgsfield affiliate links)	07:32	"I don't think anyone's actually using this" (Firefly Video); needed a 30-minute Adobe demo to understand the UX
Adobe	Tech & Design — same prompt, 4 models (Feb 2026)	00:38	Firefly "the worst one", no audio
Runway	Creative AI Show — Aleph 2.0 (May 2026, paying user)	02:06	"so clunky", crashing, "blew through like half my credits"
Runway	Curious Refuge — Aleph 2 vs Seedance (Jun 2026)	10:14	Start-frame editing "really compelling"; "not as good as Seedance 2.0"
Runway	CyberJungle — Omni vs Seedance 2 vs Aleph 2 (May 2026; Artist-sponsored segment 11:12–13:39)	03:27	"face consistency is a big problem"; "by far the slowest"

No independent video covers Firefly Enterprise or GenStudio; paid ambassador and Runway-sponsored videos are excluded.

Reddit, X, LinkedIn by name

- **Adobe:** a GenStudio buyer — "Don't buy GenStudio for performance marketing", "couldn't understand what we would actually end up buying" (Reddit); a Fortune-100 employee — "adobe is soo embedded" (Reddit); the objection we'll hear — "most enterprise users are already in the Adobe ecosystem" (Reddit); a designer — the only tool "that doesn't require a disclaimer conversation" (LinkedIn); an AI filmmaker's working stack — Midjourney, Kling, Higgsfield Cinema Studio — isn't in Firefly (X).
- **Runway:** four 15-s 1080p videos "burned everything" on Pro (Reddit); an annual plan bought for 1080p, 1080p gone a week later (Reddit); "flat-out lied to" over 1080p and Unlimited (X); "the editing workflow is the moat, not the generation quality" (X).

Switching signals

- **Away from Adobe:** "needs the ability to render like Higgsfield" (G2); advice to pick Higgsfield or Freepik over Firefly for value (Reddit).
- **Away from Runway:** "Time to find a new video generation platform" after Unlimited closed (Reddit).
- **Toward Runway, away from Higgsfield:** a buyer chose Runway because Higgsfield's Unlimited covered only Seedance 1.5 (Reddit); others left Higgsfield over throttled "unlimited" and bot support (Reddit).
- No brand publicly says "we moved from X to Y because..." — the signals are creators and small teams.

4C — Commercial motion

Pricing teardown

	Adobe	Runway
Tiers	Firefly \$9.99–199.99 a month; teams \$20–200 per licence (plans)	\$12–76 a month billed yearly (\$15–95 monthly); Team \$55 per seat, 2–9 seats (pricing)
Metric	Credits per user, no rollover. Enterprise: shared credits "per 1,000 Operations" on a closed rate card, only under an enterprise licence; admins "can't set a usage limit" per user (Adobe)	Credits; shared pool on Team; 1-month rollover on Max and Team; API at \$0.01 per credit (API)
What a 5-s clip really costs	Firefly Video at 1080p: \$2.50 on Pro	Gen-4.5: \$0.48–1.15 by plan
Behind "Contact sales"	Custom Models, Foundry, Creative Production, GenStudio	Enterprise: custom credits, unlimited members, SOC 2, legal indemnification; model-weight licensing
Legal tier	IP indemnity: \$10K on teams, greater of \$3M or 2× fees on enterprise (pricing)	Public enterprise terms, IP indemnity outside the cap (terms)
Discounts	Near-continuous promos, 30–50% off; the current one is US-only with a 12-month commitment, and its "unlimited" excludes Teams and Enterprise (terms)	–20% yearly; at least four price moves in a year that "drew criticism" (TechCrunch); Unlimited closed to new buyers (existing plans run to Nov 30, 2026); 1080p removed from Unlimited after annual sales
Ladder gap	—	A full 9-seat Team is \$5,940 a year; the mid-market AE is asked to "turn \$30k starts into \$100k+" (job post) — no published step in between

Sales hiring

	Role	Region	Seniority and pay	Signal
Adobe	Mgr, Enterprise Sales — Firefly (R170909)	San Jose, SF, NY, Lehi	First-line manager of 5+ specialists; \$258,000–421,900	A Firefly specialist overlay on the installed base
Adobe	Product Specialist — Firefly (R170840)	New York	7+ years; \$263,300–381,350	Cross-sells with the account directors
Runway	VP / Director, Enterprise Revenue (post)	US remote	\$400–475K OTE; 3–4 managers × 3–5 pods	An enterprise machine sized for growth
Runway	Enterprise AE (post) · Mid-Market AE (post)	North America · London	\$275–340K OTE · 40–60 accounts, commission held until consumption milestones	Paid on usage, not on signature
Runway	Head of Sales Development, BDR, GTM Engineer (post)	US	The first SDR leader, for a team of "currently two people"	PLG to sales-assist, being built now — 9 of 45 open roles are commercial (Ashby)

Read-out. Runway is at a stage close to Higgsfield's, one step ahead: it is building an enterprise sales machine right now — AEs in the US and Europe, pods paid on usage, a first SDR team, product signals routed to sales. Adobe is the mature player: the enterprise channel already exists, so Firefly is sold as a cross-sell on top of Creative Cloud by specialists working with the existing account teams. Its job is to defend and expand its base, not to build a channel.

Public sales collateral

- **Adobe case studies:** Too Faced (Estée Lauder) campaign in 16 days ([Adobe](#)); Newell Brands — Paper Mate packaging 75% faster with Custom Models ([Adobe](#)); IBM — 26× engagement vs its paid-social average ([guide](#), p. 15).
- **Adobe talks by SEs and sales:** Emily Palmer, Principal Solutions Consultant, Summit 2026 ([session](#)); Johan Lopes Helgesson, Firefly for business webinar ([webinar](#)); EMEA Firefly sales leaders pitch the same CMO / CIO / CFO story on LinkedIn ([Alexa Churchman-Mountbatten](#); [Sabine Fix](#)).
- **Runway case studies:** Bonjour moved to Enterprise after "hitting a wall on credits" ([Runway](#)); Chime, Miro and 8 more ([customers](#)).
- **Runway talks and sales posts:** Runway AI Summit — "From Pilot to Production: Creative AI Inside the Enterprise" ([summit](#)); Eddy Lee's talk, behind sign-in ([video](#)); the enterprise sales lead posts "\$40M in net new ARR" in half a quarter ([LinkedIn](#)) and a first \$1M quarter for one rep ([LinkedIn](#)).
- **Procurement shortcut:** Runway sells through the OpenAI Marketplace ([Runway](#)) — Higgsfield and Adobe are listed there too.

4D — Synthesis

Runway — teardown

- **Real strength:** an enterprise machine on a creator base — \$200M ARR ([Bloomberg](#)), NRR above 300% by its CRO's account, public IP indemnity, SOC 2 Type II and ISO 27001, forward-deployed creatives — and independent studios already building on it.
- **Real weakness:** its own model doesn't lead (#27 of 48), and customers don't trust its meter — Unlimited closed, 1080p removed after annual sales, misses billed.
- **The complaint we win on: paying for misses.** "Even failed generations still use up paid account credits" ([G2](#)). Our answer: the cost of every generation is shown before it runs, and packs come at fixed rates (1.3, 1.5). Honest check: our own reviews mention generations "stuck on processing, leading to wasted points" ([Trustpilot](#)) — so this has to be fixed in the enterprise terms, not claimed.
- **Watch:** Runway Ads — product feed → ads → results on Meta, Google and TikTok. Adweek reports it live ([Adweek](#)), Runway's page still says early access. It meets our ChatGPT /marketing flow head-on.

Adobe — teardown

- **Real strength:** already on every desk, and legal signs off — 20,000 enterprises, Photoshop and Premiere as the finishing layer, Firefly Foundry for brand models.
- **Real weakness:** the moat and the models people want live in different places. Users call partner models "definitely the best part of firefly" ([TR](#)), yet Seedance and Kling are kept off business plans and the indemnity covers only Adobe's own and GA Google and OpenAI models.
- **The complaint we win on: "I can't tell what I'm buying."** "couldn't understand what we would actually end up buying" ([Reddit](#)). Our answer: one page — seats by role, credit packs, every model (1.5).
- **Watch:** the two are partly one camp — Gen-4.5 inside Firefly, Adobe Ventures in Runway, Runway inside Premiere (4A).

The exploitable gap, tied to 1.2. Teams already make AI video outside Adobe's business plans — one AI filmmaker's working stack of Midjourney, Kling and Higgsfield isn't in Firefly ([X](#)) — and Runway's buyers don't trust its meter. The people who make CPG content already use Higgsfield: job ads at Reckitt and at 8 of 20 agency networks, Monks and McCann among them, name it (Part 5). We don't win on logos or models; we win by putting that existing usage under one readable contract, with the production stack and the team that make it pay off.

Where we don't fight yet: certifications and indemnity text — Runway shows SOC 2 Type II, ISO 27001 and public indemnity, Higgsfield says "SOC 2-aligned" and "contract-backed" ([Enterprise](#)); brand-model training (1.3); hero-film quality.

Must be true before the first enterprise call: a written answer on where Seedance and Kling prompts go — Adobe's own disclosure will be quoted back at us ([Adobe](#)).

1.2 says	Part 4 evidence	Verdict
Adobe is on every desk	20,000 enterprises; the finishing layer; Firefly sold as a cross-sell on the installed base	Supports
Runway is the AI-native standard for studios	Silverside builds on its API; its visible base is small studios and creators	Supports — it meets us at the land stage
Our people already use us	Job ads name Higgsfield (Part 5); switching posts name Higgsfield	Supports, per account — verified on the first call
A production stack, not just models	Models and agents are table stakes at both	Supports — the stack and its scale are the difference
A team that knows how to drive it	Runway also sells forward-deployed creatives	Partly — our edge is scale: 150+ in-house creatives (CEO on 20VC) and the Academy
Seedance and Kling are an edge over Adobe	Adobe keeps them off business plans because prompts go to their China-based developers	Partly contradicts — an edge only with a written data-routing answer

Evidence: one coder, manual themes; review bases are self-selected and mostly incentivized. Full review registers and quotes are in the support files of this folder.

Part 5 — Surprise us

The conquest map: Higgsfield is already in the agencies' job ads — not yet in the brands'

I checked creative job ads at 50 large CPG makers and 20 agency networks for three tools — 108 unique ads, up to Oct 6, 2026:

Employers naming the tool	Adobe Firefly	Runway	Higgsfield
50 CPG brands	14 (28%)	8 (16%)	1 (2%) — Reckitt
20 agency networks	15 (75%)	16 (80%)	8 (40%)

Read: brands hire for Firefly; agencies hire for Runway and Firefly — and already for Higgsfield. Our users sit inside the agencies that serve CPG, not yet inside the brands. That is the land-and-expand path from 1.1 and the two rivals from 1.2.

Agency networks whose job ads name Higgsfield: VML, McCann, FCB, Accenture Song / Droga5, Havas, Monks, OLIVER / Brandtech+, Artefact 3000.

Three plays

1. Agency networks — make it a standard.

- Signal: 8 of 20 networks name Higgsfield next to Runway and Firefly — Monks lists it among "essentials" for its AI artists ([Monks](#)), FP7 McCann for motion and AI art direction ([LinkedIn](#)).
- Play: sell the head of production a governed standard for what their people already use — a workspace per client, no training on client content, one bill (1.5).
- First line: *"Your job ads ask for Higgsfield. Let's make it a standard, not a personal card."*

2. Adobe brands — come in through Adobe, not around it.

- Signal: Firefly appears at 14 of 50 brands — P&G alone has 8 ads; Nestlé, PepsiCo, Unilever and Estée Lauder are on the list.
- Play: Seedance and Kling inside Premiere and Photoshop through our plugin, under enterprise terms — Adobe offers both to individuals only ([Adobe](#)).
- First line: *"Your team lives in Premiere. Two of the strongest video models aren't on your Adobe plan — our plugin puts them there, under your enterprise terms."*

3. Runway accounts — one brief, both tools.

- Signal: Runway is in the job ads of 16 of 20 agency networks — and in all 8 that also name Higgsfield.
- Play: a bake-off on the client's own brief, same reviewers, scored on cost per approved asset.
- First line: *"Your AI artists already use both. Let's settle it on one of your briefs."*

Why it works: competitor deals are signed at the holding or brand level; tools are chosen by the people who make the work. We win the second and grow into the first.

How to run it weekly: search job boards for three tool names → tag each brand and agency by stack → route it to its play → first email the same week.

Method: a fixed sample, not a revenue top-50; an employer counts once it has at least one verified ad naming the tool. This is hiring demand, not market share or usage — and "not found" does not mean "not used". 67 of the 108 ads have no confirmed publication date.

5.1 — Expert check: what one conversation changed

I tested my GTM assumptions with Sergey — 25+ years in marketing, in senior roles at SAP, Nestlé, Danone and Evyap.

What he told me	What I changed
His main point: enterprise isn't sold top-down — a company comes to you because one person inside saw it work	Land bottom-up; the enterprise deal turns that usage into a contract → riskiest assumption (Part 0), play 1 (Part 5)
The CMO owns the portfolio and P&L and is far from ads; content is decided one level below	Buyer = head of content production (Part 0); no CMOs on the target list (1.4)
Agencies are losing the concept business; the ones that survive adapt one idea across 100–500 touchpoints	Land through agencies and studios (1.1) — job ads confirm it: 8 of 20 networks already hire for Higgsfield (Part 5)
Big brands run content hubs where one or two people produce for 150 markets	Price the few who produce, not headcount: seats plus credit packs (1.5)
What matters is how fast a million formats come out, with how few people, and whether every one holds the brand book — a corporation can't afford the off-brand frame a blogger gets away with	Lead with speed: one approved creative into hundreds of on-brand variants, in days vs. weeks (1.1); no brand model is a gap (1.3)
"Show me why I can't do it myself" — digital teams already wire models together	The first rival is DIY: we win on the production stack, not the model list (1.2, 1.3)

One expert conversation: hypotheses, not proof. Next step — test them with the owner of a live content-production workflow.